

WB-XIC, Lab7:

Wyjaśnienia konwolucyjnych sieci neuronowych c.d.

Hubert Baniecki

hbaniecki@gmail.com | <http://hbaniecki.com>

Outline

1. SHAP
2. Neuron & layer explanations
3. Check!
4. Concept-based explanations (they exist)
5. Global explanations (intuition)
6. Project

LIME for image: superpixels and image segmentation



Label: standard poodle
Probability: 0.18
Explanation Fit: 0.37



Label: goose
Probability: 0.15
Explanation Fit: 0.55



Saliency maps (*vanilla* gradients)

The recipe for this approach is:

1. Perform a forward pass of the image of interest.
2. Compute the gradient of class score of interest with respect to the input pixels:

$$E_{grad}(I_0) = \frac{\delta S_c}{\delta I} \Big|_{I=I_0}$$


Here we set all other classes to zero.

3. Visualize the gradients. You can either show the absolute values or highlight negative and positive contributions separately.


Smoothgrad: average multiple explanations for an image with added noise

Grad-Cam: gradient explanation tailored to CNN (ReLU, last Conv2d)


Soup Bowl (vanilla)




Soup Bowl (Smoothgrad)



Soup Bowl (Grad-Cam)




Shapley values: game theory



Shapley values: math

$$\phi_j(val) = \sum_{S \subseteq \{1, \dots, p\} \setminus \{j\}} \frac{|S|! (p - |S| - 1)!}{p!} (val(S \cup \{j\}) - val(S))$$

SHapley Additive exPlanations (SHAP)



Definition 1 Additive feature attribution methods have an explanation model that is a linear function of binary variables:

$$g(z') = \phi_0 + \sum_{i=1}^M \phi_i z'_i, \quad (1)$$

where $z' \in \{0, 1\}^M$, M is the number of simplified input features, and $\phi_i \in \mathbb{R}$.

S. M. Lundberg & S. Lee. **A Unified Approach to Interpreting Model Predictions**. *NeurIPS*, 2017.

<https://dl.acm.org/doi/10.5555/3295222.3295230>

SHAP

1. (model-agnostic) **KernelSHAP**: LIME + SHAP kernel
2. TreeSHAP: fast SHAP values for tree-ensemble models
3. Gradient: SHAP based on IG and Smoothgrad
4. *SHAP based on DeepLIFT <https://arxiv.org/abs/1704.02685>

Google Colab

Neuron & layer explanations

- Neuron Attribution

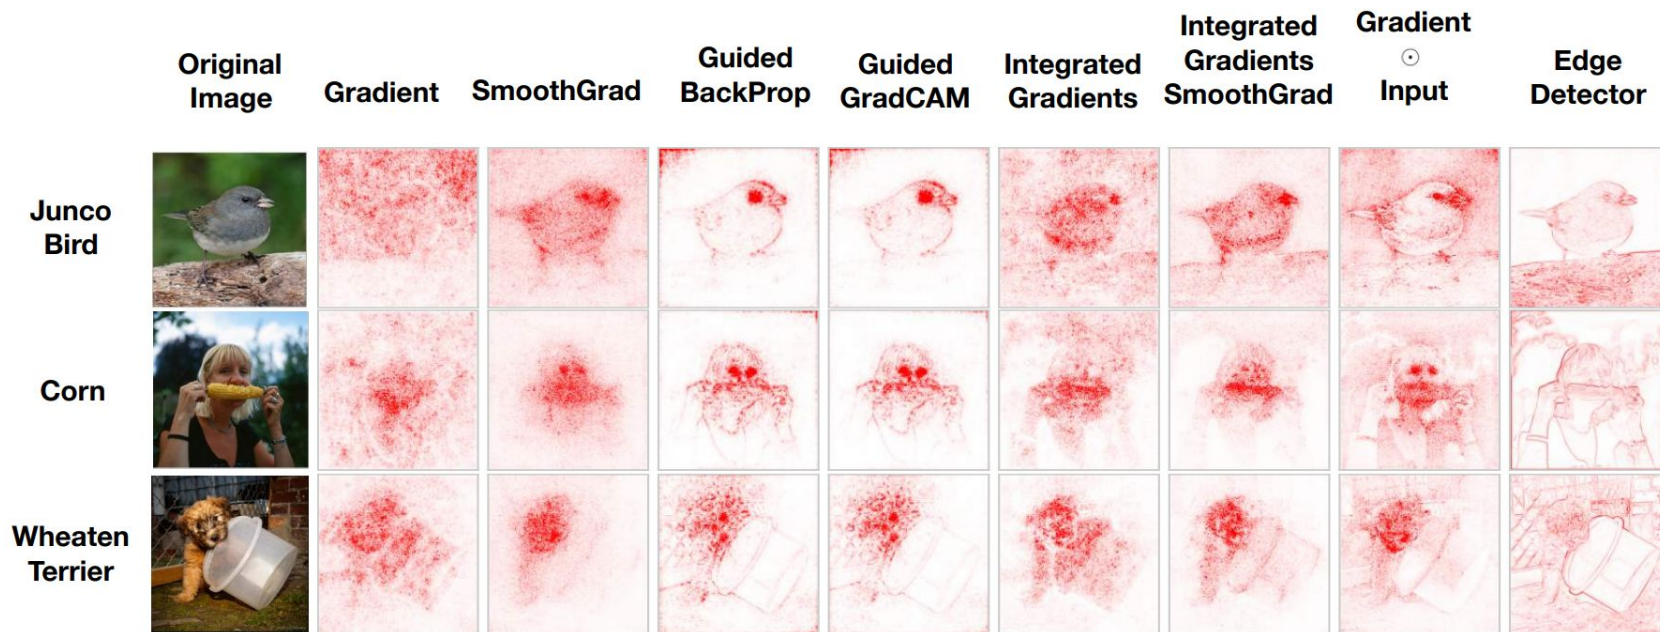
- Neuron Gradient
- Neuron Integrated Gradients
- Neuron Conductance
- Neuron DeepLift
- Neuron DeepLiftShap
- Neuron GradientShap
- Neuron Guided Backprop
- Neuron Deconvolution
- Neuron Feature Ablation

- Layer Attribution

- Layer Conductance
- Layer Activation
- Internal Influence
- Layer Gradient X Activation
- GradCAM
- Layer DeepLift
- Layer DeepLiftShap
- Layer GradientShap
- Layer Integrated Gradients
- Layer Feature Ablation
- Layer LRP

Google Colab (tasks)

Check!



J. Adebayo et al. **Sanity Checks for Saliency Maps.**
NeurIPS, 2018. <https://arxiv.org/abs/1810.03292>

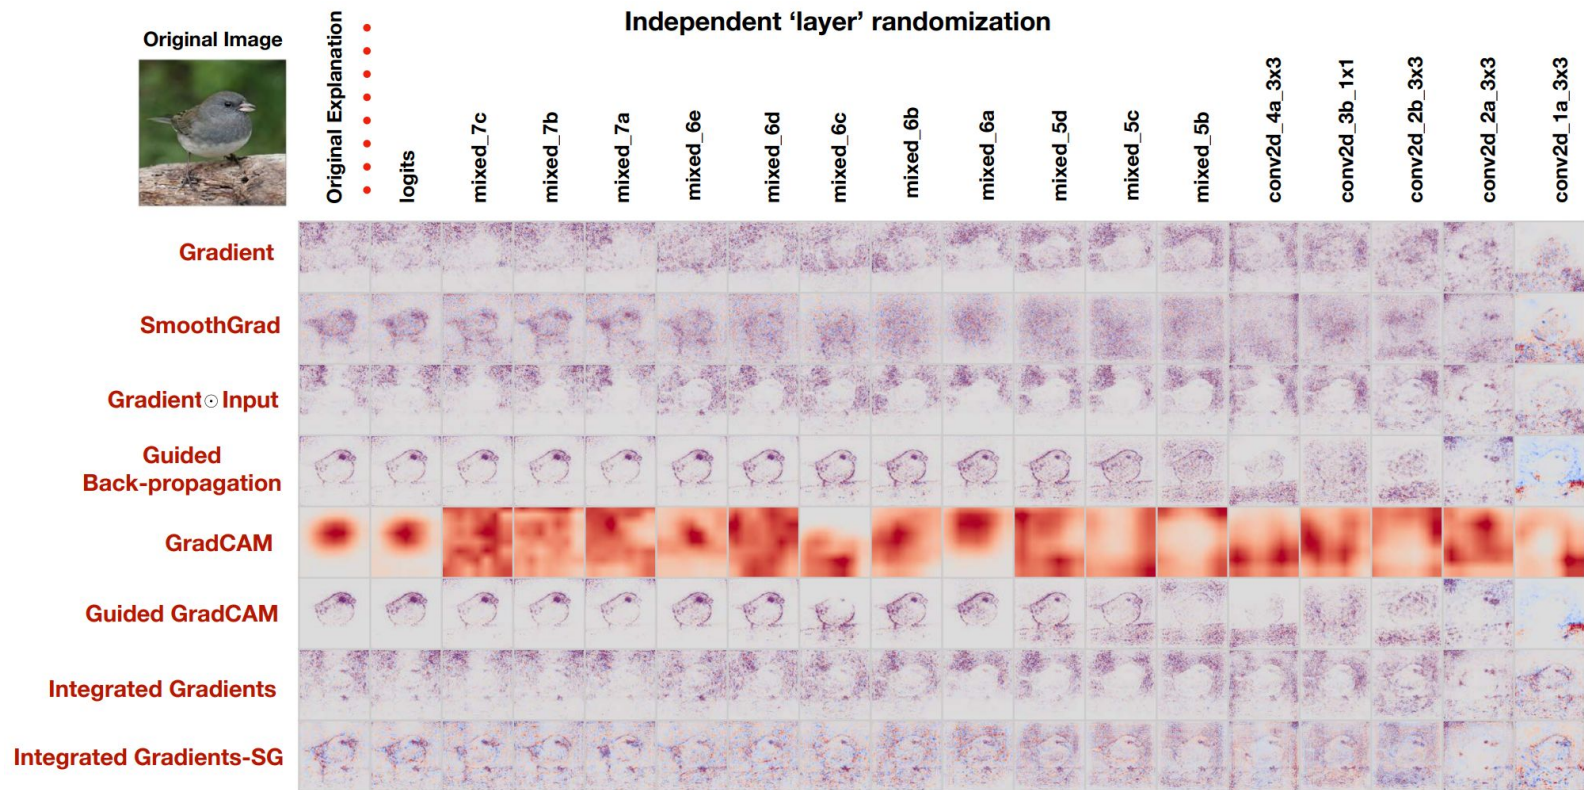


Figure 3: **Independent randomization on Inception v3 (ImageNet).** Similar to Figure 2, however

Concept-based explanations

<https://christophm.github.io/interpretable-ml-book/detecting-concepts>

- TCAV
- ConceptInterpreter
- Concept <https://captum.ai/api>
- Classifier

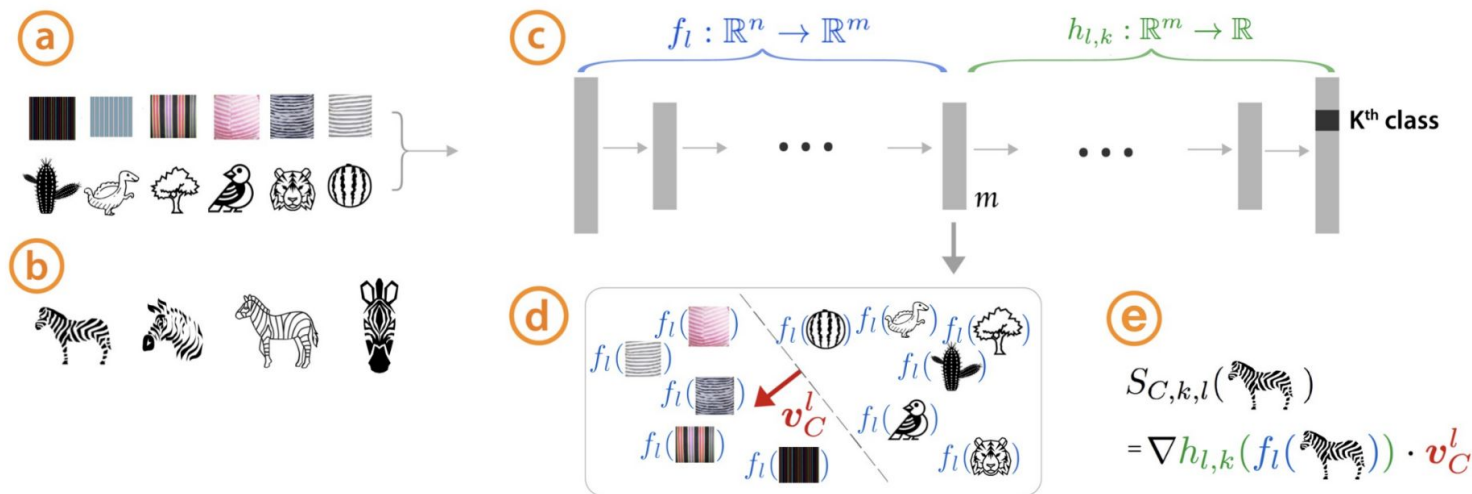


Figure 1. Testing with Concept Activation Vectors: Given a user-defined set of examples for a concept (e.g., ‘striped’), and random

B. Kim et al. **Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors (TCAV)**. ICML, 2018. <https://arxiv.org/abs/1711.11279>

Global explanations?


More formally, let $S_c(I)$ be the score of the class c , computed by the classification layer of the ConvNet for an image I . We would like to find an L_2 -regularised image, such that the score S_c is high:

$$\arg \max_I S_c(I) - \lambda \|I\|_2^2, \quad (1)$$


where λ is the regularisation parameter. A locally-optimal I can be found by the back-propagation



bell pepper



lemon



husky

Global explanations

Formally, we may pose the activation maximization problem for a unit with index j on a layer l of a network Φ as finding an image \mathbf{x}^* where:

$$\mathbf{x}^* = \arg \max_{\mathbf{x}} (\Phi_{l,j}(\mathbf{x}) - R_{\theta}(\mathbf{x}))$$

Here, $R_{\theta}(\mathbf{x})$ is a parameterized regularization: could include multiple regularizers (i.e. prior) which penalizes the search in a different way.



A. Nguyen et al. **Multifaceted Feature Visualization: Uncovering the Different Types of Features**

Learned By Each Neuron in Deep Neural Networks. /ICML 2016 <https://arxiv.org/abs/1603.03616>

Figure 4. Global explanations of the difference facets of a neuron that detects bell peppers. Diverse facets include a single, red bell pepper on a white



Figure 5. Visualizing the different facets of a neuron that detects images in the “fishing reel” class. Diverse facets include reels on

A. Nguyen et al. **Multifaceted Feature Visualization: Uncovering the Different Types of Features Learned By Each Neuron in Deep Neural Networks**. *ICML*, 2016. <https://arxiv.org/abs/1602.03616>

Global explanations: from neuron to layers

A. Nguyen et al. **Multifaceted Feature Visualization:
Uncovering the Different Types of Features Learned By
Each Neuron in Deep Neural Networks.** *ICML*, 2016.

<https://arxiv.org/abs/1602.03616>



More recently: another aggregate/cluster approach

M. Ibrahim et al. **Global Explanations of Neural Networks: Mapping the Landscape of Predictions**. *AIES*, 2019. <https://arxiv.org/abs/1902.02384>

Algorithm 1: Generating Global Attributions (GAM)

Input: local attributions

Output: medoids and corresponding members

```
/* 1. Normalize the set of local
   attributions */
foreach local attribution do
    normalized = abs(attribution) /
                sum(abs(attribution))
end

/* 2. Compute pair-wise rank
   distance matrix */
distances = []

foreach attribution1 in normalizedAttributions do
    foreach attribution2 in normalizedAttributions
        do
            distances += rankDistance(attribution1,
                                      attribution2)
        end
    end
end

/* 3. Cluster Attributions */
initialMedoids = random.choice(attributions)
for x iterations do
    foreach cluster do
        foreach attribution in cluster do
            tempMedoid = attribution;
            cost = sum(distance(attribution,
                               tempMedoid));
            reassign medoid to attribution
                minimizing cost;
        end
        update cluster membership by assigning to
            closest medoid
    end
end
```

Projekt